

Software tools for CA6000 – Local LLM and Jupyter AI

Continuing with our discussion of low code programming using AI coding assistant, we will install a large language model (LLM) and integrate it with the Jupyter AI in our Jupyter Lab on our laptop, such that we can have our own local AI coding assistant running on our laptop without the need to connect to the LLM hosted on the cloud (i.e. internet)

We will be using the following packages in the installation.

- Ollama + selected LLM
- Jupyter Lab install in a separate environment
- Jupyter AI

A. Installing Ollama (Procedure and screenshots are as of Nov 2025)

Ollama is an open-source tool that can allow us to run LLMs locally on our computer.

- Go to the Ollama download website (<https://ollama.com/download>) and choose the appropriate version for your computer to download its installer program.

Download Ollama



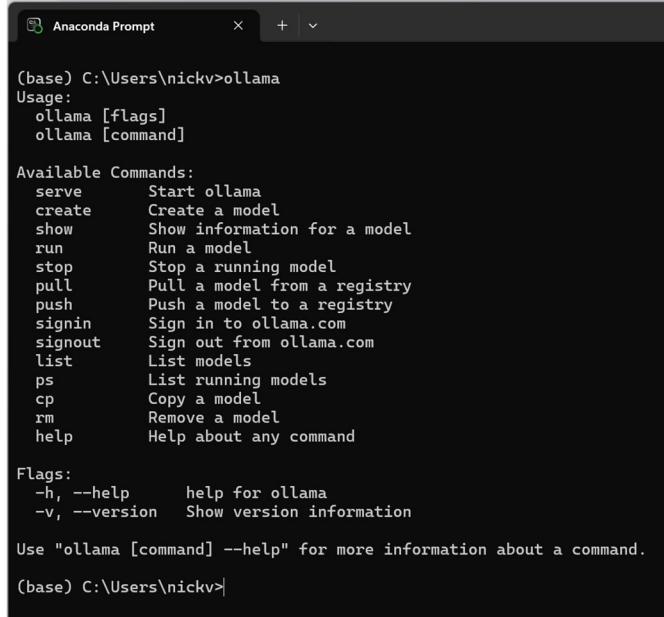
- Run the installer program - you will see a message similar to the following

Extracting files...
C:\Users\nickv\AppData\Local\Programs\Ollama\lib\ollama\cuda_v12\cublasLt64_12.dll

- Following the on-screen instruction to complete the installation. (A small Ollama icon { } may appear at the corner of your taskbar that will launch a GUI based interface if you click on it. But we will not be using this interface)



- After the installation, the Ollama server should execute automatically - you can check that it is running by opening a terminal (such as the anaconda prompt) and type the command “**ollama**”



```
(base) C:\Users\nickv>ollama
Usage:
  ollama [flags]
  ollama [command]

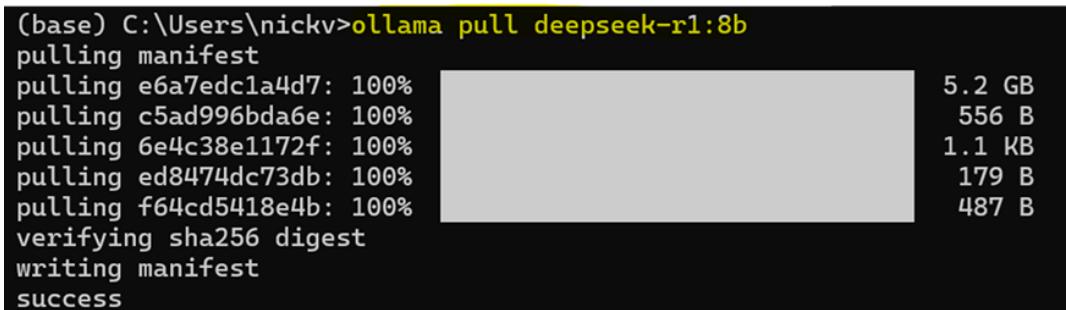
Available Commands:
  serve      Start ollama
  create     Create a model
  show       Show information for a model
  run        Run a model
  stop       Stop a running model
  pull       Pull a model from a registry
  push       Push a model to a registry
  signin    Sign in to ollama.com
  signout   Sign out from ollama.com
  list      List models
  ps        List running models
  cp        Copy a model
  rm        Remove a model
  help      Help about any command

Flags:
  -h, --help    help for ollama
  -v, --version Show version information

Use "ollama [command] --help" for more information about a command.

(base) C:\Users\nickv>
```

- You can use the “**ollama list**” command to view all the downloaded models. It should be empty at this stage as this is the first time you use Ollama.
- We will next install an LLM model to run with the Ollama. The available LLM models are listed at <https://ollama.com/library>. To download a specific model, use the “**ollama pull**” command. E.g. **ollama pull deepseek-r1:8b**



```
(base) C:\Users\nickv>ollama pull deepseek-r1:8b
pulling manifest
pulling e6a7edc1a4d7: 100% [██████████] 5.2 GB
pulling c5ad996bda6e: 100% [██████████] 556 B
pulling 6e4c38e1172f: 100% [██████████] 1.1 KB
pulling ed8474dc73db: 100% [██████████] 179 B
pulling f64cd5418e4b: 100% [██████████] 487 B
verifying sha256 digest
writing manifest
success
```

- Once the model is downloaded, we can run the model by using the “**ollama run**” command. E.g. **ollama run deepseek-r1:8b**

```
(base) C:\Users\nickv>ollama run deepseek-r1:8b
>>> |Send a message (/? for help)
```

Note: You will use this command to launch the Ollama server with the selected model in future run of your system.

- You can test run the installation by typing a message, which you should see a response such as the following . E.g., type “Hello”

```
>>> Hello
Hello! How can I assist you today? 😊

>>> Send a message (/? for help)
```

B. Optional: Installing Python Library

We can also interact with the Ollama LLM using Python code. To do so, we will install the Ollama’s Python library (which is called ‘ollama’ - a bit confusing, I know)

- Open a cmd console, or an Anaconda Prompt console and type the command ‘[pip install ollama](#)’

```
C:\Users\aschvun>pip install ollama
Collecting ollama
  Downloading ollama-0.6.0-py3-none-any.whl.metadata (4.3 kB)
Collecting httpx>=0.27 (from ollama)
  Downloading httpx-0.28.1-py3-none-any.whl.metadata (7.1 kB)
Collecting pydantic>=2.9 (from ollama)
  Downloading pydantic-2.12.0-py3-none-any.whl.metadata (92 kB)
```

The Ollama Python library provides functions such as `generate()` , `chat()` that you can use in a Python program to ‘chat’ with the LLM.

```
from ollama import generate
model = "deepseek-r1:1.5b"

response = generate(model, 'What does 30*5 equal to?')
print("DeepSeek:", response["response"])

>>> -
= RESTART: C:\Users\aschvun\AppData\Local\Programs\Python\Python311\AI6120 - DeepSeek0.py
DeepSeek: <think>
To calculate 30 multiplied by 5, I first recognize that 30 is a multiple of 10.
I then break down the multiplication into two parts: multiplying 3 by 5 and then adding a

### Final Answer:
\[[
30 \times 5 = \boxed{150}
\]
```

C. Installing Jupyter Lab in a new Environment

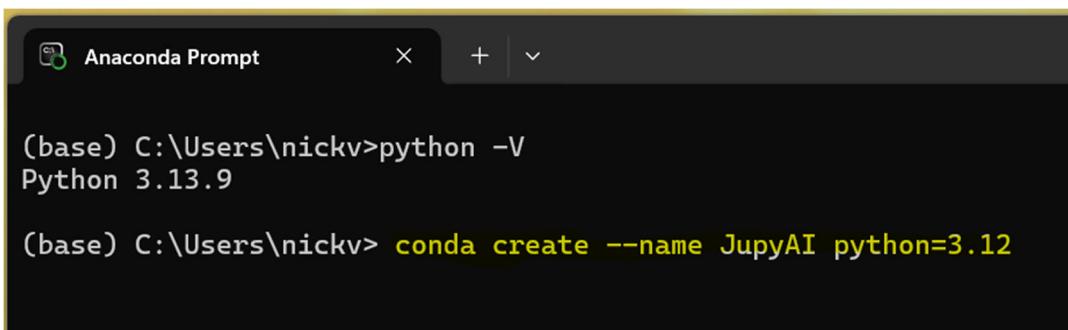
Jupyter AI is an extension package that enable the integration of generative AI capabilities directly in the Jupyter Lab (and Jupyter Notebook) platform. Once it is installed, we can use it to interact with a LLM model that is running remotely in the cloud through internet, or in our case, running locally on our computer through the Ollama.

Due to rapid changes in software versions that may not always be 100% backward compatible, it is common to have situation that a package can only work with certain version of another package (known as dependency). As such, we need to be aware of such dependency constraint when installing packages. For example, based on the latest (Nov 2025) information, Jupyter AI will only work with Python version 3.12.

Python provides the concept of ‘Environment’ where Python code can be run in a self-contained space, with its own interpreter, libraries and installed packages. So we can install multiple versions of the same software on the same computer, but with each running in its own environment. As such, it is recommended that we create a new Environment when we install the Jupyter AI with the Jupyter Lab as described below.

To begin with, we will create an Environment ‘JupyAI’ to install the Jupyter AI and Jupyter Lab (even though you had already installed the Jupyter Lab earlier. But that version will be running in a different environment – by default, the ‘base’).

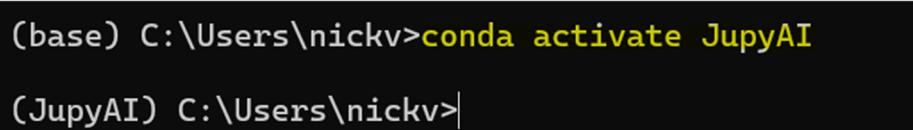
- Launch the Anaconda prompt to open a terminal (i.e. console). Check the version of Python that is already installed in our computer (and notice that it is in the environment “base”). We then create a new environment “JupyAI” and install it with Python 3.12.



```
(base) C:\Users\nickv>python -V
Python 3.13.9

(base) C:\Users\nickv> conda create --name JupyAI python=3.12
```

- Follow the installation instruction, and once it is done, we can change to the new environment by activating it. Notice that we are now in the ‘JupyAI’ environment.



```
(base) C:\Users\nickv>conda activate JupyAI

(JupyAI) C:\Users\nickv>
```

- We will next install the Jupyter lab in this JupyAI environment, using the conda install instruction and specifying the conda-forge channel.

```
(JupyAI) C:\Users\nickv>conda install -c conda-forge jupyterlab
3 channel Terms of Service accepted
Channels:
- conda-forge
- defaults
Platform: win-64
Collecting package metadata (repodata.json): done
Solving environment: done
```

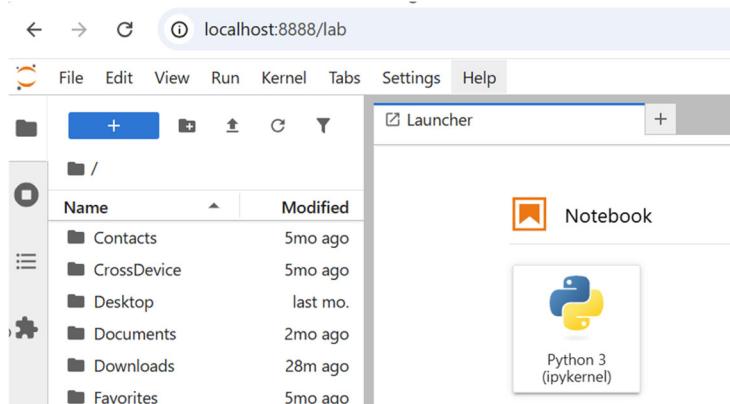
- Following the instructions and check the version that is installed

```
(JupyAI) C:\Users\nickv>jupyter lab --version
4.4.10
```

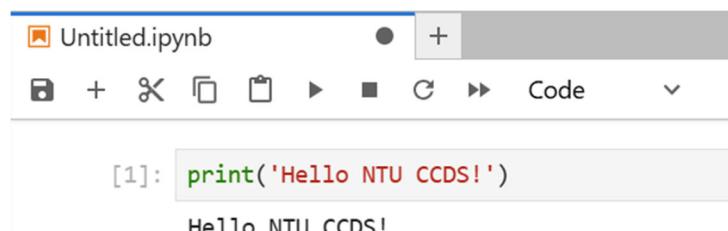
- Start the Jupyter Server program in this JupyAI environment

```
(JupyAI) C:\Users\nickv>jupyter lab
[I 2025-11-09 11:47:02.876 ServerApp] jupyter_ai | extension v
[I 2025-11-09 11:47:02.877 ServerApp] jupyter_lsp | extension v
```

- You should see that your default browser is launched with a webpage with url at localhost:8888



- Select Python 3 Notebook and run a code to check that it is working as expected



- Terminate the Jupyter Lab server and close the browser (to get ready for installing and integration of Jupyter AI)

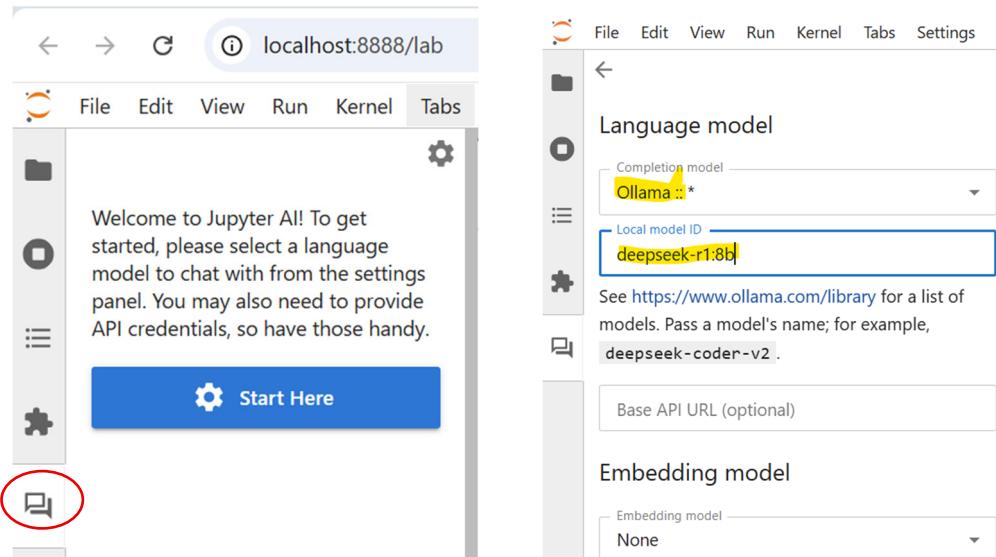
D. Installing Jupyter AI (in the new Environment)

Here are the steps to install and integrate the Jupyter AI into the Jupyter Lab.

- Use conda install instruction to install the Jupyter AI with LangChain for Ollama

```
(JupyAI) C:\Users\nickv>conda install -c conda-forge jupyter-ai langchain-ollama
3 channel Terms of Service accepted
Channels:
- conda-forge
- defaults
Platform: win-64
Collecting package metadata (repodata.json): - |
```

- Launch the Jupyter Lab and you should see the Jupyter AI Tab become available in the Tabs panel. Click the Start Here button and configure it with the elected LLM. In this case, Ollama and the earlier installed LLM deepseek-r1:8b.



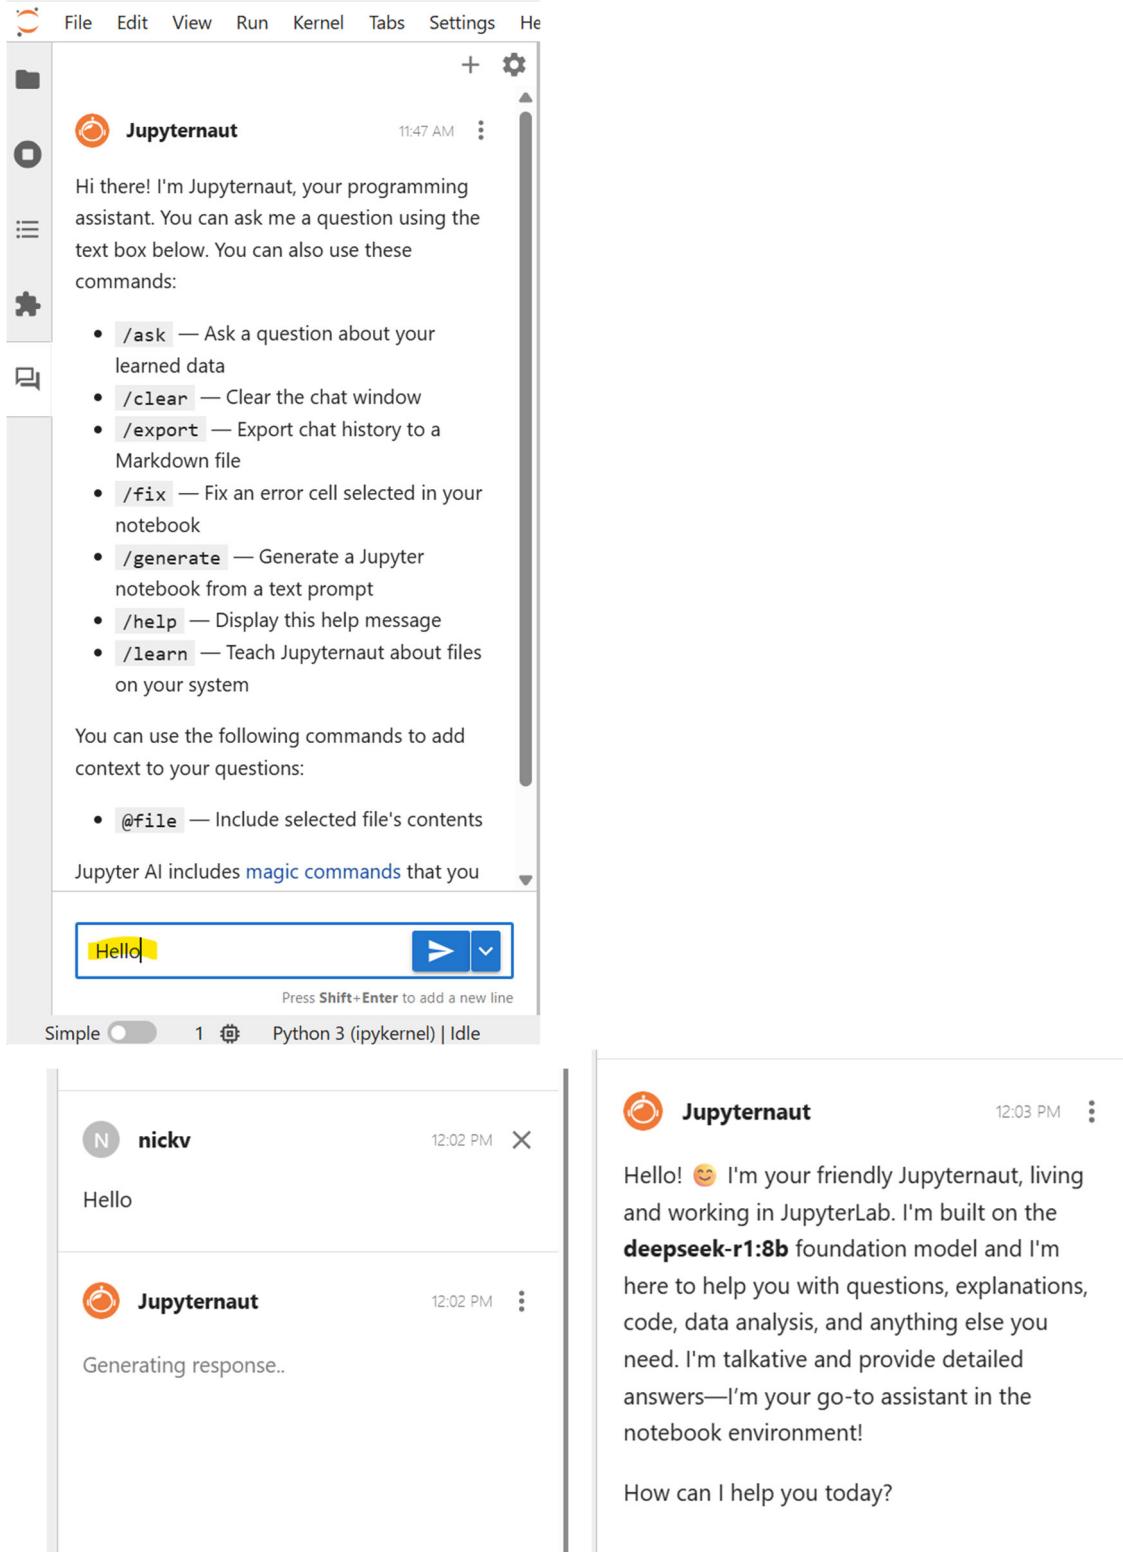
- Click Save changes button.



- Terminate the Jupyter Lab (server and browser).

Start a new session of Jupyter Lab server

- Now the Jupyter AI will display its interface (Jupyternaut) that allows you to key in messages (E.g. 'Hello').



The screenshot shows the JupyterLab interface with the Jupyternaut AI assistant. The top navigation bar includes File, Edit, View, Run, Kernel, Tabs, Settings, Help, and a user icon. On the left, there's a sidebar with icons for file operations like New, Open, Save, and a gear for settings.

Jupyternaut (11:47 AM)

Hi there! I'm Jupyternaut, your programming assistant. You can ask me a question using the text box below. You can also use these commands:

- /ask — Ask a question about your learned data
- /clear — Clear the chat window
- /export — Export chat history to a Markdown file
- /fix — Fix an error cell selected in your notebook
- /generate — Generate a Jupyter notebook from a text prompt
- /help — Display this help message
- /learn — Teach Jupyternaut about files on your system

You can use the following commands to add context to your questions:

- @file — Include selected file's contents

Jupyter AI includes [magic commands](#) that you can use.

In the bottom text input field, the user has typed "Hello". Below it, a button says "Press Shift+Enter to add a new line".

Simple 1 Python 3 (ipykernel) | Idle

nickv (12:02 PM)

Hello

Jupyternaut (12:02 PM)

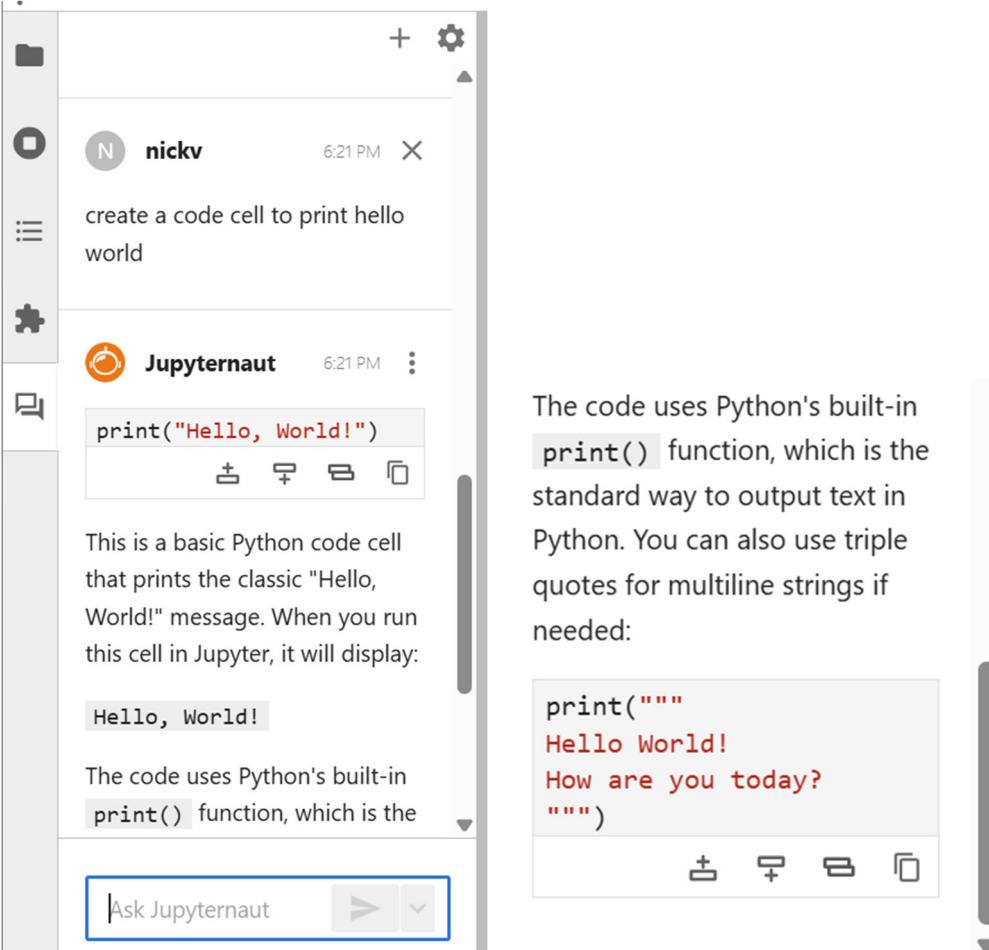
Generating response..

Jupyternaut (12:03 PM)

Hello! 😊 I'm your friendly Jupyternaut, living and working in JupyterLab. I'm built on the **deepseek-r1:8b** foundation model and I'm here to help you with questions, explanations, code, data analysis, and anything else you need. I'm talkative and provide detailed answers—I'm your go-to assistant in the notebook environment!

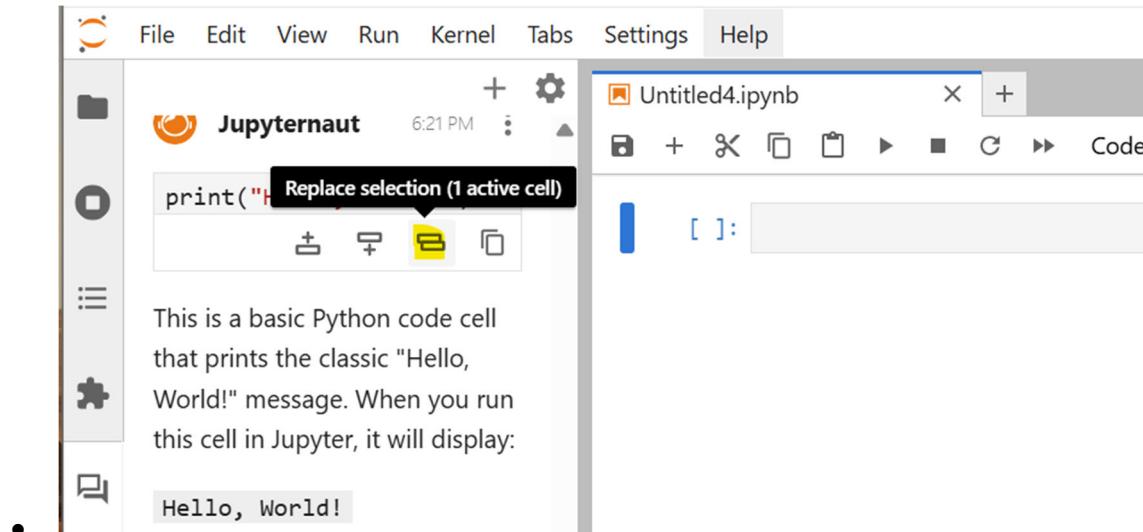
How can I help you today?

- Another example – generating a python code: “create a code cell to print hello world”

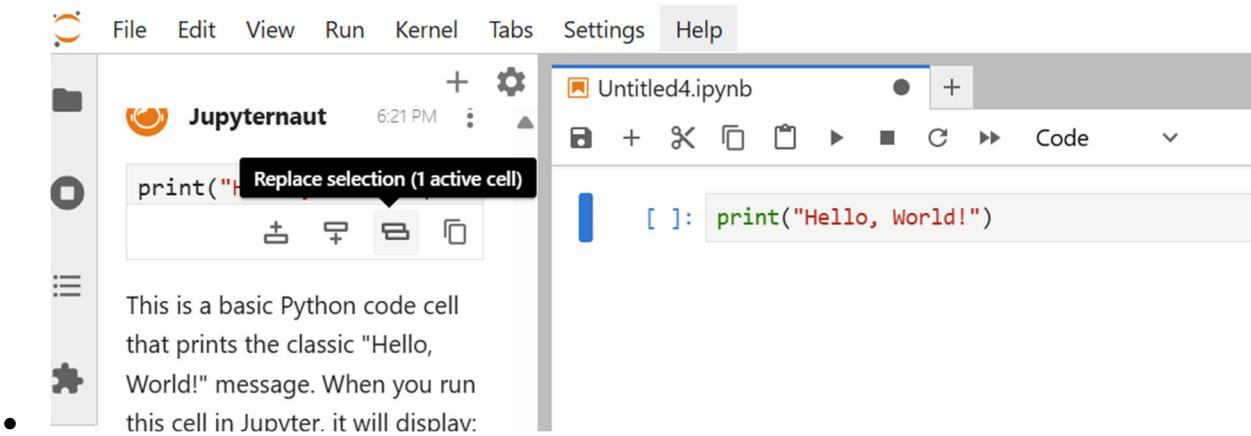


The screenshot shows the Jupyter Notebook interface. On the left, there's a sidebar with icons for file operations. The main area has two code cells. The first cell, titled "nickv" (6:21 PM), contains the text "create a code cell to print hello world". The second cell, titled "Jupyternaut" (6:21 PM), contains the Python code `print("Hello, World!")`. Below the code cell, a text block explains: "This is a basic Python code cell that prints the classic "Hello, World!" message. When you run this cell in Jupyter, it will display:". To the right of this text, the output of the code cell is shown: "Hello, World!". A note below the output says: "The code uses Python's built-in `print()` function, which is the". On the far right, another code cell is shown with the code `print("""Hello World!
How are you today?
""")`.

- The generated code can be directly copied to the active cell by clicking on the button “Replace selection”



This screenshot shows the same Jupyter Notebook interface as above, but with a mouse cursor hovering over the "Replace selection" button in the toolbar of the active cell. The active cell contains the Python code `print("Hello, World!")`. The text below the code cell is identical to the previous screenshot: "This is a basic Python code cell that prints the classic "Hello, World!" message. When you run this cell in Jupyter, it will display:". The output "Hello, World!" is also present.



The screenshot shows the Jupyter Notebook interface. The top menu bar includes File, Edit, View, Run, Kernel, Tabs, Settings, and Help. A toolbar below the menu has icons for file operations like new, open, save, and run. The main area displays a notebook titled "Untitled4.ipynb". A code cell contains the Python command `print("Hello, World!")`. A tooltip window is open over the code, showing the text: "This is a basic Python code cell that prints the classic "Hello, World!" message. When you run this cell in Jupyter, it will display:". At the bottom of the screen, there is a status bar with the text "6:21 PM".